

```
import matplotlib.pyplot as plt
import pandas as pd
import seaborn as sns
import numpy as np
import sklearn
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
from sklearn.compose import ColumnTransformer
from sklearn.preprocessing import LabelEncoder
from sklearn.preprocessing import OneHotEncoder, StandardScaler
import warnings
warnings.filterwarnings('ignore')
from sklearn.tree import DecisionTreeRegressor
from sklearn.feature_selection import RFE
from sklearn.preprocessing import MinMaxScaler, StandardScaler
import joblib
import datetime
from sklearn.feature_selection import SelectFromModel
from sklearn.model_selection import cross_val_score, GridSearchCV
from sklearn.linear_model import Ridge, Lasso

In [1]: data = pd.read_csv("exams.csv")

In [2]: data

Out [3]:
```

	gender	race/ethnicity	parental level of education	lunch	test preparation course	math score	reading score	writing score
0	female	group D	some college	standard	completed	59.0	70.0	78.0
1	male	group D	associate's degree	standard	none	96.0	93.0	87.0
2	female	group D	some college	free/reduced	none	57.0	76.0	77.0
3	male	group B	some college	free/reduced	none	70.0	70.0	63.0
4	female	group D	associate's degree	standard	none	83.0	85.0	86.0
...
995	male	group C	some college	standard	none	77.0	77.0	71.0
996	male	group C	some college	standard	none	80.0	66.0	66.0
997	female	group A	high school	standard	completed	67.0	86.0	86.0
998	male	group E	high school	standard	none	80.0	72.0	62.0
999	male	group D	high school	standard	none	58.0	47.0	45.0

1000 rows × 8 columns

```
In [4]: data.head()

Out [4]:
```

	gender	race/ethnicity	parental level of education	lunch	test preparation course	math score	reading score	writing score
0	female	group D	some college	standard	completed	59.0	70.0	78.0
1	male	group D	associate's degree	standard	none	96.0	93.0	87.0
2	female	group D	some college	free/reduced	none	57.0	76.0	77.0
3	male	group B	some college	free/reduced	none	70.0	70.0	63.0
4	female	group D	associate's degree	standard	none	83.0	85.0	86.0

```
In [5]: data.tail()

Out [5]:
```

	gender	race/ethnicity	parental level of education	lunch	test preparation course	math score	reading score	writing score
995	male	group C	some college	standard	none	77.0	77.0	71.0
996	male	group C	some college	standard	none	80.0	66.0	66.0
997	female	group A	high school	standard	completed	67.0	86.0	86.0
998	male	group E	high school	standard	none	80.0	72.0	62.0
999	male	group D	high school	standard	none	58.0	47.0	45.0

```
In [6]: data.shape

Out [6]: (1000, 8)

In [7]: data.info()
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 8 columns):
#   Column                Non-Null Count  Dtype
---  ---
0   gender                 1000 non-null   object
1   race/ethnicity         1000 non-null   object
2   parental level of education 1000 non-null   object
3   lunch                  1000 non-null   object
4   test preparation course 1000 non-null   object
5   math score             975 non-null    float64
6   reading score          993 non-null    float64
7   writing score           989 non-null    float64
dtypes: float64(3), object(5)
memory usage: 62.6+ KB

In [8]: data.describe()

Out [8]:
```

	math score	reading score	writing score
count	975.000000	993.000000	989.000000
mean	67.841026	70.379568	69.165824
std	15.210716	14.108466	14.995555
min	15.000000	25.000000	15.000000
25%	58.000000	61.000000	59.000000
50%	68.000000	70.000000	70.000000
75%	79.000000	80.000000	80.000000
max	100.000000	100.000000	100.000000

```
In [9]: data.isnull().sum()

Out [9]:
gender                0
race/ethnicity        0
parental level of education 0
lunch                 0
test preparation course 0
math score            25
reading score         11
writing score         11
dtype: int64

In [10]: data['math score'].mean()

Out [10]: 67.84102564102564

In [11]: data['math score'].fillna(data['math score'].mean(), inplace = True)

In [12]: data['reading score'].mean()

Out [12]: 70.37956760322555

In [13]: data['reading score'].fillna(data['reading score'].mean(), inplace = True)

In [14]: data['writing score'].mean()

Out [14]: 69.16582496471183

In [15]: data['writing score'].fillna(data['writing score'].mean(), inplace = True)

In [16]: data.isnull().sum()

Out [16]:
gender                0
race/ethnicity        0
parental level of education 0
lunch                 0
test preparation course 0
math score            0
reading score         0
writing score         0
dtype: int64

In [17]: mode_value = data['math score'].mode().iloc[0]
print("Mode:", mode_value)
Mode: 62.0


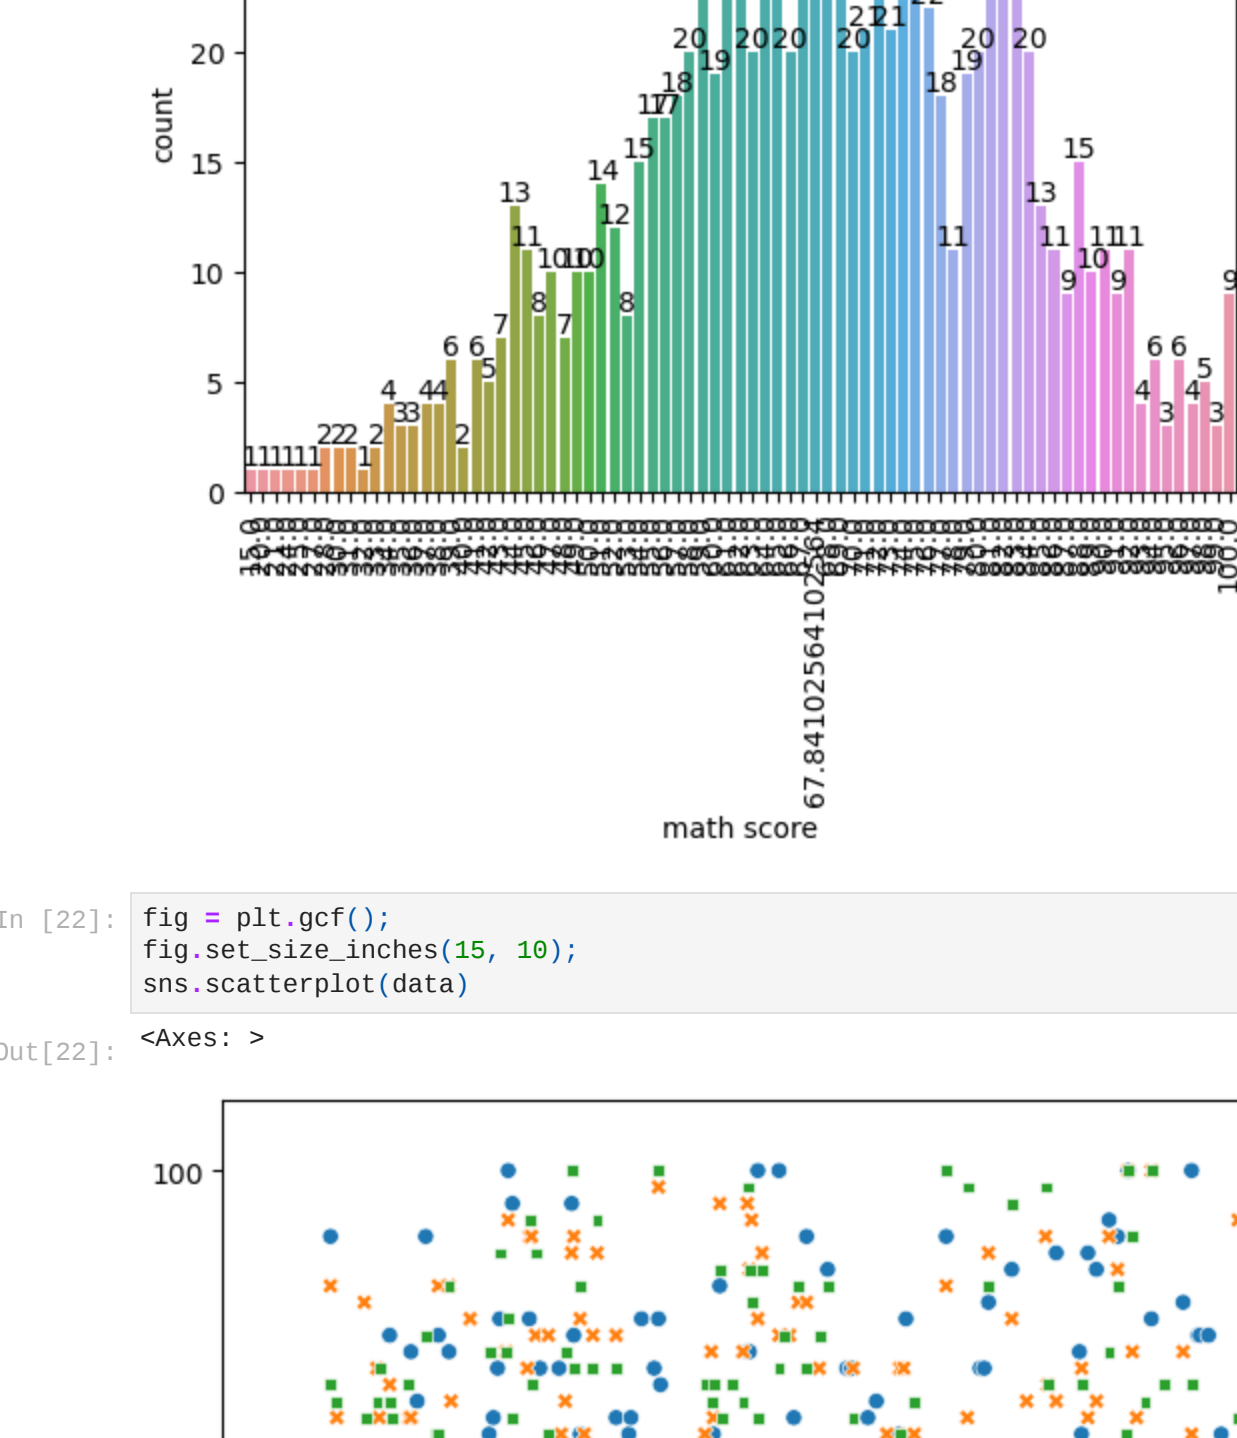
In [18]: mode_value = data['reading score'].mode().iloc[0]
print("Mode:", mode_value)
Mode: 72.0

In [19]: mode_value = data['writing score'].mode().iloc[0]
print("Mode:", mode_value)
Mode: 72.0

In [20]: sns.heatmap(data.isnull(), yticklabels=False, cmap="Paired")

Out [20]: <Axes: >


```




```
In [23]: data

Out [23]:
```

	gender	race/ethnicity	parental level of education	lunch	test preparation course	math score	reading score	writing score
0	female	group D	some college	standard	completed	59.0	70.0	78.0
1	male	group D	associate's degree	standard	none	96.0	93.0	87.0
2	female	group D	some college	free/reduced	none	57.0	76.0	77.0
3	male	group B	some college	free/reduced	none	70.0	70.0	63.0
4	female	group D	associate's degree	standard	none	83.0	85.0	86.0
...
995	male	group C	some college	standard	none	77.0	77.0	71.0
996	male	group C	some college	standard	none	80.0	66.0	66.0
997	female	group A	high school	standard	completed	67.0	86.0	86.0
998	male	group E	high school	standard	none	80.0	72.0	62.0
999	male	group D	high school	standard	none	58.0	47.0	45.0

1000 rows × 8 columns



```
In [26]: encoder = LabelEncoder()

In [27]: data['gender'] = encoder.fit_transform(data['gender'])
data['parental level of education'] = encoder.fit_transform(data['parental level of education'])
data['lunch'] = encoder.fit_transform(data['lunch'])

In [28]: data.shape

Out [28]: (1000, 8)

In [29]: categorical_col = ['gender', 'parental level of education', 'lunch']
encoder = OneHotEncoder(drop='first', sparse=False)
encoder_cols = pd.DataFrame(encoder.fit_transform(data[categorical_col]).columns=encoder.get_feature_names_out(categorical_col))

In [30]: encoder_cols

Out [30]:
```

	gender_1	parental level of education_1	parental level of education_2	parental level of education_3	parental level of education_4	parental level of education_5	lunch_1
0	0.0	0.0	0.0	0.0	0.0	1.0	0.0
1	1.0	0.0	0.0	0.0	0.0	0.0	1.0
2	0.0	0.0	0.0	0.0	1.0	0.0	0.0
3	1.0	0.0	0.0	0.0	1.0	0.0	0.0
4	0.0	0.0	0.0	0.0	0.0	0.0	1.0
...
995	1.0	0.0	0.0	0.0	1.0	0.0	1.0
996	1.0	0.0	0.0	0.0	1.0	0.0	1.0
997	0.0	0.0	1.0	0.0	0.0	0.0	1.0
998	1.0	0.0	1.0	0.0	0.0	0.0	1.0
999	1.0	0.0	1.0	0.0	0.0	0.0	1.0

1000 rows × 7 columns

```
In [31]: data

Out [31]:
```

	gender	race/ethnicity	parental level of education	lunch	test preparation course	math score	reading score	writing score
0	0	group D	4	1	completed	59.0	70.0	78.0
1	1	group D	0	1	none	96.0	93.0	87.0
2	0	group D	4	0	none	57.0	76.0	77.0
3	1	group B	4	0	none	70.0	70.0	63.0
4	0	group D	0	1	none	83.0	85.0	86.0
...
995	1	group C	4	1	none	77.0	77.0	71.0
996	1	group C	4	1	none	80.0	66.0	66.0
997	0	group A	2	1	completed	67.0	86.0	86.0
998	1	group E	2	1	none	80.0	72.0	62.0
999	1	group D	2	1	none	58.0	47.0	45.0

1000 rows × 8 columns

```
In [32]: numerical_col = ['math score', 'reading score', 'writing score']
Scaled = StandardScaler()
Scaled= pd.DataFrame(Scaled.fit_transform(data[numerical_col]), columns=numerical_col)

In [33]: Scaled

Out [33]:
```

	math score	reading score	writing score
0	-0.509943	-0.027017	0.922528
1	1.875905	1.609714	1.196179
2	-0.722173	0.399956	0.525456
3	0.143820	-0.027017	-0.418956
4	1.009812	1.040416	1.129107
...
995	0.610123	0.471118	0.123022
996	-0.809968	-0.311666	-0.212339
997	-0.056025	1.111578	1.129107
998	0.809968	0.115307	-0.480628
999	-0.659558	-1.663749	-1.620857

1000 rows × 3 columns

```
In [34]: x = pd.concat([encoder_cols, Scaled], axis=1)
y = data['gender']

In [35]: x.shape

Out [35]: (1000, 10)

In [36]: y

Out [36]:
```

	Y
0	0
1	1
2	0
3	1
4	0
...	...
995	1
996	1
997	0
998	1
999	1

gender, Length: 1000, dtype: int32

```
In [37]: X_train, X_test, Y_train, Y_test=train_test_split(x, y, test_size=0.2, random_state=42)

In [38]: model=LinearRegression()
model.fit(X_train, Y_train)
y_pred=model.predict(X_test)

In [39]: model = LinearRegression()
model.fit(X_train, Y_train)
y_pred = model.predict(X_test)

In [40]: from sklearn.metrics import r2_score, mean_squared_error

In [41]: r2= r2_score(Y_test, y_pred)

In [42]: mae = mean_absolute_error(Y_test, y_pred)
mse = mean_squared_error(Y_test, y_pred)
rmse = np.sqrt(mse)
r2 = r2_score(Y_test, y_pred)

In [43]: print("Mean Absolute Error", mae)
print("Mean Squared Error", mse)
print("Root Mean Squared Error", rmse)
print("R2 Score", r2)
Mean Absolute Error 7.8441464399222440e-15
Mean Squared error 9.9062701161537030e-15
Root Mean Squared Error 9.9062701161537030e-15
R2 Score 1.0

In [ ]: lr_model.fit(X_train, Y_train)
lr_prediction = lr_model.predict(X_test)
lr_mae = mean_absolute_error(Y_test, lr_prediction)
lr_mse = mean_squared_error(Y_test, lr_prediction)
lr_r2 = r2_score(Y_test, lr_prediction)

In [ ]: print("Linear MAE", lr_mae)
print("Linear MSE", lr_mse)
print("Linear R2", lr_r2)

In [44]: lr_model = LinearRegression()
lr_scores = cross_val_score(lr_model, X_train, Y_train, cv=5)

In [45]: ridge_model= Ridge(alpha=1.0)
ridge_scores = cross_val_score(ridge_model, X_train, Y_train, cv = 5)

In [46]: Lasso_model= Lasso(alpha=1.0)
Lasso_scores = cross_val_score(Lasso_model, X_train, Y_train, cv = 5)

In [47]: Lasso_model.fit(X_train, Y_train)
Lasso_prediction = Lasso_model.predict(X_test)
Lasso_mae = mean_absolute_error(Y_test, Lasso_prediction)
Lasso_mse = mean_squared_error(Y_test, Lasso_prediction)
Lasso_r2 = r2_score(Y_test, Lasso_prediction)

In [48]: print("Lasso MAE", Lasso_mae)
print("Lasso MSE", Lasso_mse)
print("Lasso R2", Lasso_r2)
Lasso MAE 0.49981250909090905
Lasso MSE 0.24982656250000005
Lasso R2 -0.001810776942356096

In [49]: ridge_model.fit(X_train, Y_train)
ridge_prediction = ridge_model.predict(X_test)
ridge_mae = mean_absolute_error(Y_test, ridge_prediction)
ridge_mse = mean_squared_error(Y_test, ridge_prediction)
ridge_r2 = r2_score(Y_test, ridge_prediction)

In [50]: print("Lasso MAE", ridge_mae)
print("Lasso MSE", ridge_mse)
print("Lasso R2", ridge_r2)
Lasso MAE 0.002791417359187404
Lasso MSE 1.1679960764055893e-05
Lasso R2 0.9999531632811467

In [51]: from sklearn.linear_model import HuberRegressor
scaler=StandardScaler()
X_scaled = scaler.fit_transform(X_scaled)
huber = HuberRegressor(epsilon=1.35)
huber.fit(X_scaled, Y_test)
huber_prediction = huber.predict(X_scaled)
huber_mae = mean_absolute_error(Y_test, huber_prediction)
huber_mse = mean_squared_error(Y_test, huber_prediction)
huber_r2 = r2_score(Y_test, huber_prediction)
print("huber mae:", huber_mae)
print("huber mse:", huber_mse)
print("huber r2:", huber_r2)
huber mae: 1.837563237255328e-11
huber mse: 5.549402194085089e-22
huber rse: 2.3527169722049530e-11
huber r2: 1.0

In [ ]:
```