

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error
from sklearn.preprocessing import LabelEncoder, OneHotEncoder, StandardScaler
import warnings
warnings.filterwarnings('ignore')
from sklearn.tree import DecisionTreeRegressor
from sklearn.feature_selection import RFE
from sklearn.preprocessing import MinMaxScaler, StandardScaler
import joblib
import datetime
from sklearn.feature_selection import SelectFromModel
from sklearn.model_selection import cross_val_score, GridSearchCV
from sklearn.linear_model import Ridge, Lasso
```

```
In [2]: df = pd.read_csv('mushrooms.csv')
```

```
In [3]: df
```

Out[3]:

	class	cap- shape	cap- surface	cap- color	bruises	odor	gill- attachment	gill- spacing	gill- size	gill- color	...	stalk- surface- below- ring
0	p	x	s	n	t	p	f	c	n	k	...	s
1	e	x	s	y	t	a	f	c	b	k	...	s
2	e	b	s	w	t	l	f	c	b	n	...	s
3	p	x	y	w	t	p	f	c	n	n	...	s
4	e	x	s	g	f	n	f	w	b	k	...	s
...
8119	e	k	s	n	f	n	a	c	b	y	...	s
8120	e	x	s	n	f	n	a	c	b	y	...	s
8121	e	f	s	n	f	n	a	c	b	n	...	s
8122	p	k	y	n	f	y	f	c	n	b	...	k
8123	e	x	s	n	f	n	a	c	b	y	...	s

8124 rows × 23 columns



In [4]:

df.head()

Out[4]:

	class	cap-shape	cap-surface	cap-color	bruises	odor	gill-attachment	gill-spacing	gill-size	gill-color	...	stalk-surface-below-ring	
0	p	x	s	n	t	p	f	c	n	k	...	s	
1	e	x	s	y	t	a	f	c	b	k	...	s	
2	e	b	s	w	t	l	f	c	b	n	...	s	
3	p	x	y	w	t	p	f	c	n	n	...	s	
4	e	x	s	g	f	n	f	w	b	k	...	s	

5 rows × 23 columns

In [5]:

df.tail()

Out[5]:

	class	cap-shape	cap-surface	cap-color	bruises	odor	gill-attachment	gill-spacing	gill-size	gill-color	...	stalk-surface-below-ring	
8119	e	k	s	n	f	n	a	c	b	y	...	s	
8120	e	x	s	n	f	n	a	c	b	y	...	s	
8121	e	f	s	n	f	n	a	c	b	n	...	s	
8122	p	k	y	n	f	y	f	c	n	b	...	k	
8123	e	x	s	n	f	n	a	c	b	y	...	s	


5 rows × 23 columns

```
In [6]: df.isnull().sum()
```

```
Out[6]: class                0
        cap-shape            0
        cap-surface          0
        cap-color            0
        bruises              0
        odor                 0
        gill-attachment       0
        gill-spacing          0
        gill-size             0
        gill-color            0
        stalk-shape           0
        stalk-root            0
        stalk-surface-above-ring 0
        stalk-surface-below-ring 0
        stalk-color-above-ring 0
        stalk-color-below-ring 0
        veil-type             0
        veil-color            0
        ring-number           0
        ring-type             0
        spore-print-color      0
        population            0
        habitat               0
        dtype: int64
```

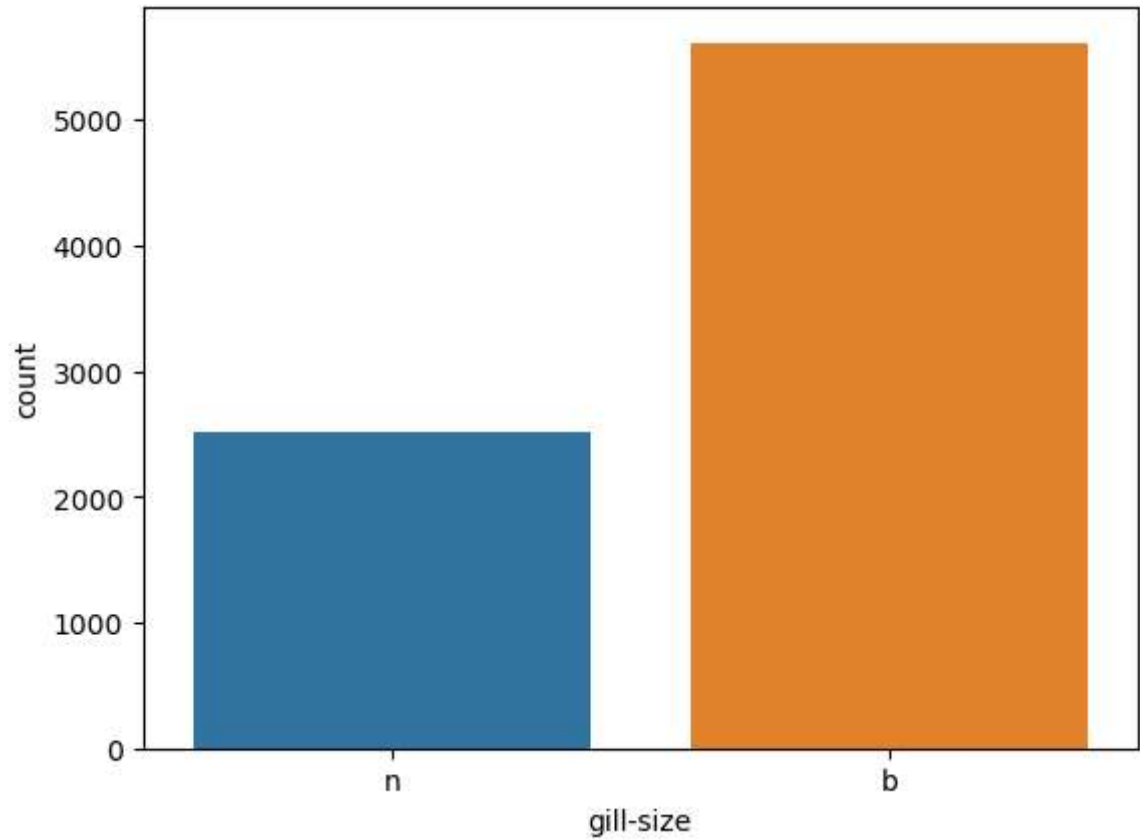
```
In [7]: sns.heatmap(df.isnull(),yticklabels=False,cmap="Paired")
```

```
Out[7]: <Axes: >
```




PERFORMING EDA

```
In [10]: sns.countplot(x=df['gill-size'])  
plt.show()
```




```
In [12]: plt.hist(df['cap-shape'], bins=10, color='skyblue')  
  
plt.xlabel('population')  
plt.ylabel('Y-Axis Label')  
plt.title('cap-shape')  
plt.show()
```



```
In [18]: sns.countplot(x='class', data=df, palette='viridis')
plt.xlabel('class')
plt.ylabel('Count')
plt.title('Count Plot')

plt.show()
```

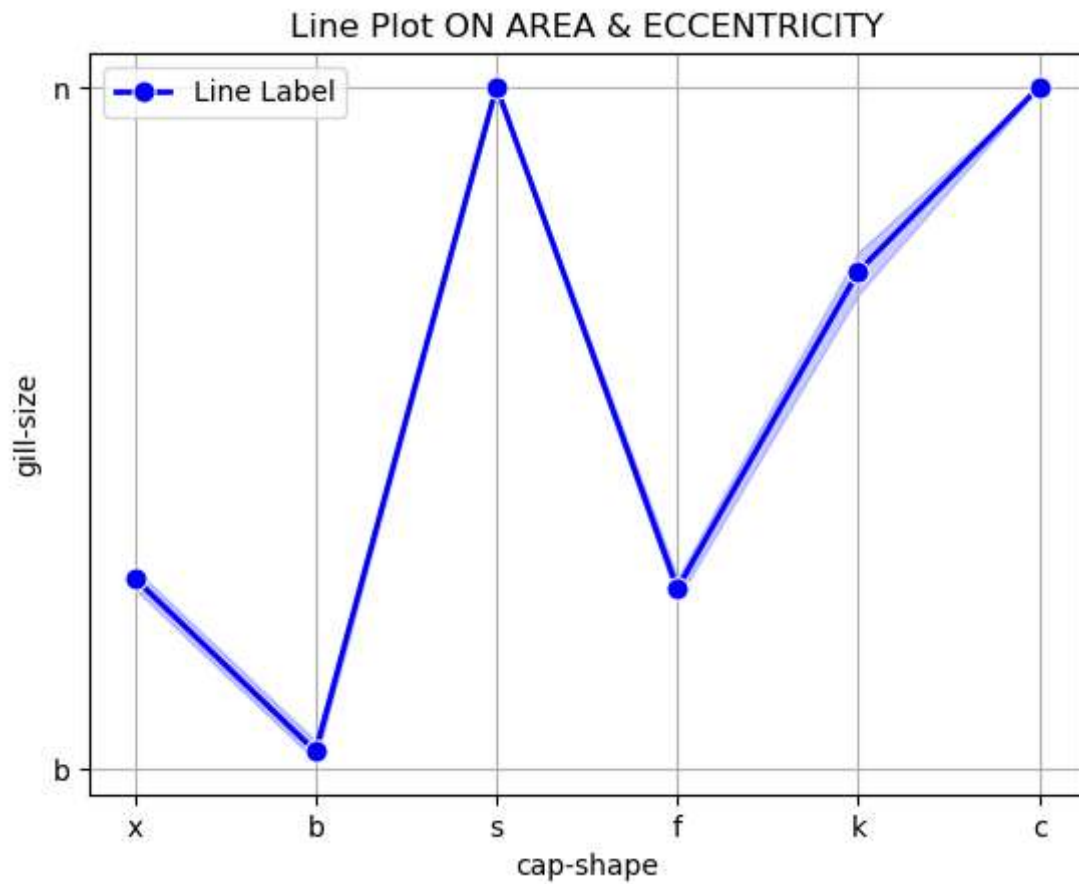


```
In [21]: df_area=df['cap-shape'].value_counts()
```

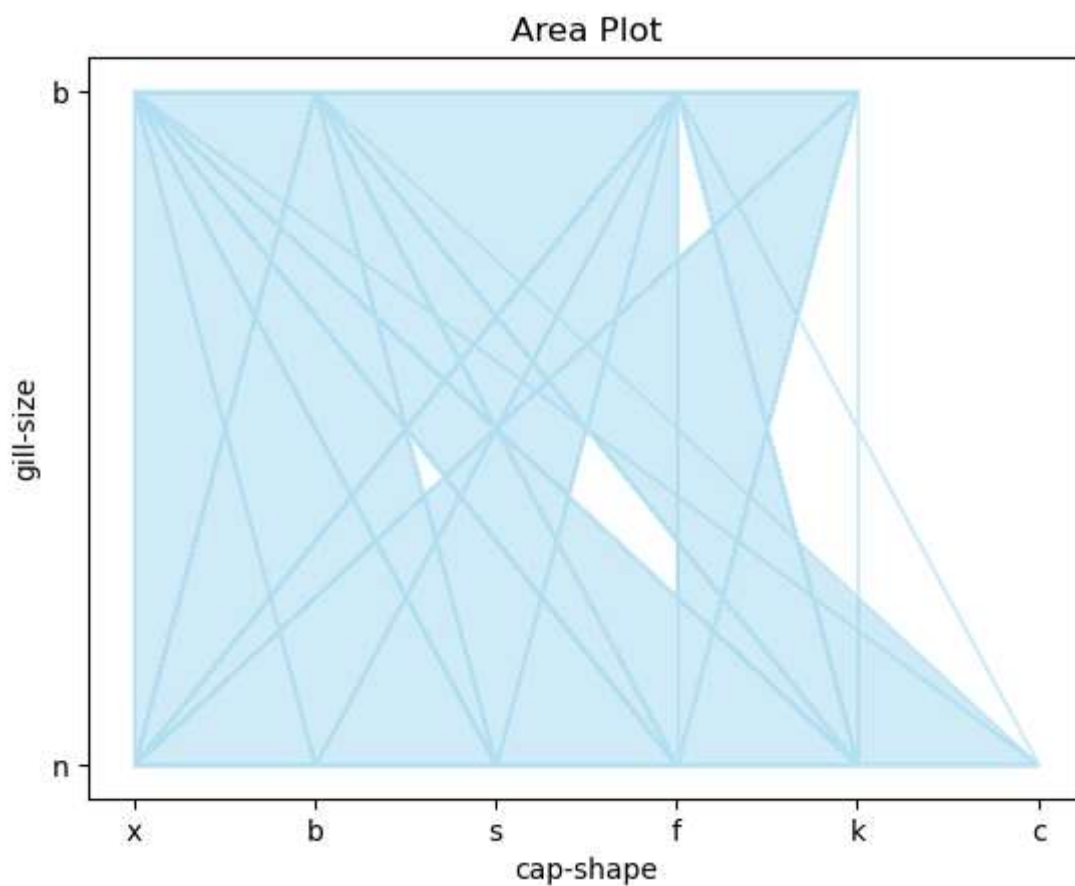
```
In [22]: df_area
```

```
Out[22]: x    3656
f    3152
k     828
b     452
s      32
c       4
Name: cap-shape, dtype: int64
```

```
In [23]: sns.lineplot(x='cap-shape', y='gill-size', data=df, marker='o', color='b', lin  
plt.xlabel('cap-shape')  
plt.ylabel('gill-size')  
plt.title('Line Plot ON AREA & ECCENTRICITY')  
plt.legend()  
plt.grid(True)  
plt.show()
```




```
In [27]: plt.fill_between(df['cap-shape'], df['gill-size'], color='skyblue', alpha=0.4)
plt.xlabel('cap-shape')
plt.ylabel('gill-size')
plt.title('Area Plot')
plt.show()
```




```
In [31]: df_pie = df['cap-shape'].value_counts()
```


```
In [32]: df_pie
```

```
Out[32]: x    3656
         f    3152
         k     828
         b     452
         s      32
         c       4
         Name: cap-shape, dtype: int64
```

```
In [36]: plt.pie(df_pie,labels=df_pie.index,autopct="%.2f%%",radius =1.2,colors=['skybl  
plt.show()
```




```
In [42]: sns.countplot(x=df['cap-shape'],hue=df['gill-size'])  
plt.show()
```



```
In [44]: sns.stripplot(x='cap-shape',y='gill-size',data=df)
```

```
Out[44]: <Axes: xlabel='cap-shape', ylabel='gill-size'>
```



```
In [48]: cat_data=df.select_dtypes(include=object)
num_data=df.select_dtypes(exclude=object)
```

```
In [49]: num_data
```

```
Out[49]:
```

```
0
1
2
3
4
...
8119
8120
8121
8122
8123
```

8124 rows × 0 columns

In [50]: cat_data

Out[50]:

	class	cap-shape	cap-surface	cap-color	bruises	odor	gill-attachment	gill-spacing	gill-size	gill-color	...	stalk-surface-below-ring
0	p	x	s	n	t	p	f	c	n	k	...	s
1	e	x	s	y	t	a	f	c	b	k	...	s
2	e	b	s	w	t	l	f	c	b	n	...	s
3	p	x	y	w	t	p	f	c	n	n	...	s
4	e	x	s	g	f	n	f	w	b	k	...	s
...
8119	e	k	s	n	f	n	a	c	b	y	...	s
8120	e	x	s	n	f	n	a	c	b	y	...	s
8121	e	f	s	n	f	n	a	c	b	n	...	s
8122	p	k	y	n	f	y	f	c	n	b	...	k
8123	e	x	s	n	f	n	a	c	b	y	...	s

8124 rows × 23 columns



In [52]: df.dtypes

```
Out[52]: class                object
cap-shape                object
cap-surface              object
cap-color                object
bruises                  object
odor                    object
gill-attachment          object
gill-spacing             object
gill-size                object
gill-color               object
stalk-shape              object
stalk-root               object
stalk-surface-above-ring  object
stalk-surface-below-ring  object
stalk-color-above-ring    object
stalk-color-below-ring    object
veil-type                object
veil-color               object
ring-number              object
ring-type                object
spore-print-color         object
population               object
habitat                  object
dtype: object
```

In [53]: df

Out[53]:


	class	cap-shape	cap-surface	cap-color	bruises	odor	gill-attachment	gill-spacing	gill-size	gill-color	...	stalk-surface-below-ring
0	p	x	s	n	t	p	f	c	n	k	...	s
1	e	x	s	y	t	a	f	c	b	k	...	s
2	e	b	s	w	t	l	f	c	b	n	...	s
3	p	x	y	w	t	p	f	c	n	n	...	s
4	e	x	s	g	f	n	f	w	b	k	...	s
...
8119	e	k	s	n	f	n	a	c	b	y	...	s
8120	e	x	s	n	f	n	a	c	b	y	...	s
8121	e	f	s	n	f	n	a	c	b	n	...	s
8122	p	k	y	n	f	y	f	c	n	b	...	k
8123	e	x	s	n	f	n	a	c	b	y	...	s

8124 rows × 23 columns



In [54]: encoder=LabelEncoder()

In [55]: categorical_col = ['class', 'cap-shape', 'cap-surface', 'cap-color', 'bruises', 'od
encoder = OneHotEncoder(drop='first', sparse=False)
encoder_cols = pd.DataFrame(encoder.fit_transform(df[categorical_col]), columns




In [56]: encoder_cols

Out[56]:

	class_p	cap-shape_c	cap-shape_f	cap-shape_k	cap-shape_s	cap-shape_x	cap-surface_g	cap-surface_s	cap-surface_y
0	1.0	0.0	0.0	0.0	0.0	1.0	0.0	1.0	0.0
1	0.0	0.0	0.0	0.0	0.0	1.0	0.0	1.0	0.0
2	0.0	0.0	0.0	0.0	0.0	0.0	0.0	1.0	0.0
3	1.0	0.0	0.0	0.0	0.0	1.0	0.0	0.0	1.0
4	0.0	0.0	0.0	0.0	0.0	1.0	0.0	1.0	0.0
...
8119	0.0	0.0	0.0	1.0	0.0	0.0	0.0	1.0	0.0
8120	0.0	0.0	0.0	0.0	0.0	1.0	0.0	1.0	0.0
8121	0.0	0.0	1.0	0.0	0.0	0.0	0.0	1.0	0.0
8122	1.0	0.0	0.0	1.0	0.0	0.0	0.0	0.0	1.0
8123	0.0	0.0	0.0	0.0	0.0	1.0	0.0	1.0	0.0

8124 rows × 88 columns




In [64]: x = pd.concat([encoder_cols],axis=1)
y = df['class']

In [65]: x

Out[65]:

	class_p	cap-shape_c	cap-shape_f	cap-shape_k	cap-shape_s	cap-shape_x	cap-surface_g	cap-surface_s	cap-surface_y
0	1.0	0.0	0.0	0.0	0.0	1.0	0.0	1.0	0.0
1	0.0	0.0	0.0	0.0	0.0	1.0	0.0	1.0	0.0
2	0.0	0.0	0.0	0.0	0.0	0.0	0.0	1.0	0.0
3	1.0	0.0	0.0	0.0	0.0	1.0	0.0	0.0	1.0
4	0.0	0.0	0.0	0.0	0.0	1.0	0.0	1.0	0.0
...
8119	0.0	0.0	0.0	1.0	0.0	0.0	0.0	1.0	0.0
8120	0.0	0.0	0.0	0.0	0.0	1.0	0.0	1.0	0.0
8121	0.0	0.0	1.0	0.0	0.0	0.0	0.0	1.0	0.0
8122	1.0	0.0	0.0	1.0	0.0	0.0	0.0	0.0	1.0
8123	0.0	0.0	0.0	0.0	0.0	1.0	0.0	1.0	0.0

8124 rows × 88 columns



```
In [67]: y
```

```
Out[67]: 0      p
          1      e
          2      e
          3      p
          4      e
          ..
        8119     e
        8120     e
        8121     e
        8122     p
        8123     e
        Name: class, Length: 8124, dtype: object
```

```
In [68]: x_train,x_test,y_train,y_test= train_test_split(x,y,test_size=0.2,random_state
```

```
In [69]: from sklearn.preprocessing import StandardScaler
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn import datasets
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error
from sklearn.preprocessing import LabelEncoder,OneHotEncoder,StandardScaler
import warnings
from sklearn.svm import SVC
from sklearn.metrics import accuracy_score, classification_report
from sklearn.metrics import confusion_matrix
```

```
In [70]: scaler = StandardScaler()
```

```
In [71]: x_train = scaler.fit_transform(x_train)
x_test = scaler.fit_transform(x_test)
```

```
In [72]: svc = SVC(kernel='linear')
```

```
In [73]: svc.fit(x_train,y_train)
```

```
Out[73]: SVC(kernel='linear')
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [74]: y_pred = svc.predict(x_test)
```



```
In [75]: y_pred
```

```
Out[75]: array(['e', 'p', 'p', ..., 'p', 'p', 'p'], dtype=object)
```

```
In [76]: acc = accuracy_score(y_test,y_pred)
print('Accuracy:{:.2f}%'. format(acc*100))
```

Accuracy:100.00%

```
In [77]: print(classification_report(y_test,y_pred))
```


	precision	recall	f1-score	support
e	1.00	1.00	1.00	843
p	1.00	1.00	1.00	782
accuracy			1.00	1625
macro avg	1.00	1.00	1.00	1625
weighted avg	1.00	1.00	1.00	1625

```
In [78]: cm = confusion_matrix(y_test,y_pred)
print('confusion matrix:')
print(cm)
```

confusion matrix:
[[843 0]
[0 782]]

```
In [79]: sns.heatmap(cm,annot=True)
```

```
Out[79]: <Axes: >
```



```
In [80]: from sklearn.metrics import roc_curve, auc
from sklearn.preprocessing import label_binarize
from sklearn.multiclass import OneVsRestClassifier
```

```
In [81]: yb = label_binarize(y,classes=[0,1,2])
```

```
In [82]: nc = yb.shape[1]
```

```
In [83]: classifier = OneVsRestClassifier(SVC(kernel='linear',probability=True,random_s
```

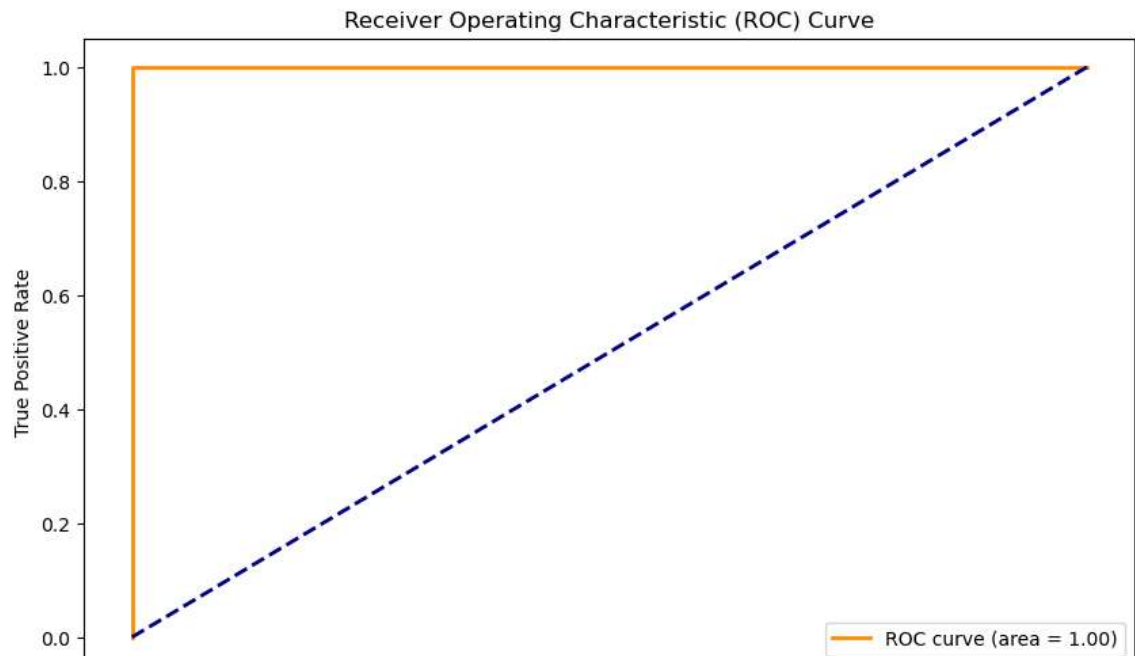
```
In [84]: y_score = classifier.fit(x_train,y_train).decision_function(x_test)
```

```
In [85]: y_score_2d = y_score.reshape(-1, 1)
```

```
In [87]: y_test_binary = (y_test == 'p').astype(int)
fpr, tpr, _ = roc_curve(y_test_binary, y_score)
```

```
In [88]: roc_auc = auc(fpr, tpr)
```

```
In [89]: plt.figure(figsize=(10, 6))
plt.plot(fpr, tpr, color='darkorange', lw=2, label=f'ROC curve (area = {roc_auc})')
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve')
plt.legend(loc='lower right')
plt.show()
In
```



```
In [90]: from sklearn.svm import SVC
from sklearn.model_selection import GridSearchCV
param_grid = {
    'C': [0.1, 1, 10, 100],
    'kernel': ['linear'],
}
svm = SVC()
grid_search = GridSearchCV(svm, param_grid, cv=5, n_jobs=-1)
grid_search.fit(x_train, y_train)
print("Best hyperparameters found: ", grid_search.best_params_)
print("Best accuracy on the validation set: {:.2f}".format(grid_search.best_score_))
```

Best hyperparameters found: {'C': 0.1, 'kernel': 'linear'}
 Best accuracy on the validation set: 1.00

```
In [91]: grid_search = GridSearchCV(svm, param_grid, cv=5, n_jobs=-1)
grid_search.fit(x_train, y_train)
```

```
Out[91]: GridSearchCV(cv=5, estimator=SVC(), n_jobs=-1,
                      param_grid={'C': [0.1, 1, 10, 100], 'kernel': ['linear']})
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [92]: print("Best hyperparameters found: ", grid_search.best_params_)
print("Best accuracy on the validation set: {:.2f}".format(grid_search.best_sc
```

```
Best hyperparameters found: {'C': 0.1, 'kernel': 'linear'}
Best accuracy on the validation set: 1.00
```

```
In [93]: best_svm = grid_search.best_estimator_
best_svm.fit(x_train, y_train)
```

```
Out[93]: SVC(C=0.1, kernel='linear')
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [94]: test_accuracy = best_svm.score(x_test, y_test)
print("Test accuracy: {:.2f}".format(test_accuracy))
```

```
Test accuracy: 1.00
```

```
In [95]: print("Test accuracy: {:.2f}".format(test_accuracy))
```

```
Test accuracy: 1.00
```

```
In [96]: from sklearn.model_selection import RandomizedSearchCV
from scipy.stats import uniform
```

```
In [97]: svm=SVC(kernel='linear')
param_dist={
    'C':uniform(loc=0,scale=10),
    'gamma':['scale','auto']+list(uniform(loc=0,scale=1).rvs(10)),
}
```

```
In [98]: n_iter_search=20
random_search = RandomizedSearchCV(svc, param_distributions=param_dist, n_iter=
random_search.fit(x_train,y_train)
```

```
Out[98]: RandomizedSearchCV(cv=5, estimator=SVC(kernel='linear'), n_iter=20, n_jobs=-
1,
          param_distributions={'C': <scipy.stats._distn_infrastructure.rv_continuous_frozen object at 0x0000020E8231F890>,
                              'gamma': ['scale', 'auto',
                                         0.6084981184617801,
                                         0.1527775985145413,
                                         0.8245404219416449,
                                         0.7976862097920863,
                                         0.6878790088079055,
                                         0.7418297071837214,
                                         0.04358746495667076,
                                         0.644865554378094,
                                         0.058261500366677876,
                                         0.1667991441025477]}},
          random_state=42)
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [99]: best_param = random_search.best_params_
best_model = random_search.best_estimator_
y_pred_2=best_model.predict(x_test)
```

```
In [100]: print(classification_report(y_test,y_pred))
```

	precision	recall	f1-score	support
e	1.00	1.00	1.00	843
p	1.00	1.00	1.00	782
accuracy			1.00	1625
macro avg	1.00	1.00	1.00	1625
weighted avg	1.00	1.00	1.00	1625

```
In [101]: cm= confusion_matrix(y_test,y_pred_2)
print(cm)
```

```
[[843  0]
 [ 0 782]]
```

```
In [102]: from sklearn import model_selection, naive_bayes, svm, metrics, feature_extraction
```

```
In [103]: x = pd.concat([encoder_cols],axis=1)
          y = df['class']
```

```
In [104]: x_train,x_test,y_train,y_test= train_test_split(x,y,test_size=0.2,random_state=42)
```

```
In [105]: from sklearn.preprocessing import MinMaxScaler
          scaler = MinMaxScaler()
          x_train = scaler.fit_transform(x_train)
          x_test = scaler.transform(x_test)
```

```
In [106]: bayes = naive_bayes.MultinomialNB()
```

```
In [107]: bayes.fit(x_train,y_train)
```

```
Out[107]: MultinomialNB()
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [108]: y_pred_nb=bayes.predict(x_test)
```

```
In [109]: accuracy=metrics.accuracy_score(y_test,y_pred_nb)
          accuracy
```

```
Out[109]: 0.9993846153846154
```

```
In [110]: print(metrics.classification_report(y_test, y_pred_nb))
```

	precision	recall	f1-score	support
e	1.00	1.00	1.00	843
p	1.00	1.00	1.00	782
accuracy			1.00	1625
macro avg	1.00	1.00	1.00	1625
weighted avg	1.00	1.00	1.00	1625

```
In [111]: cm=confusion_matrix(y_test,y_pred)
          cm
```

```
Out[111]: array([[843,  0],
                 [ 0, 782]], dtype=int64)
```

```
In [112]: yb=label_binarize(y, classes=[0,1,2])
nc = yb.shape[1]
```

```
In [120]: param_grid = {
'alpha': [0.1, 1, 10, 100],
'fit_prior': [True, False]
}
```

```
In [121]: bayes = naive_bayes.MultinomialNB()
grid_search = GridSearchCV(bayes, param_grid, cv=5)
grid_search.fit(x_train, y_train)
```

```
Out[121]: GridSearchCV(cv=5, estimator=MultinomialNB(),
                      param_grid={'alpha': [0.1, 1, 10, 100],
                                   'fit_prior': [True, False]})
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [122]: best_param = grid_search.best_params_
best_nb = naive_bayes.MultinomialNB(alpha = best_param['alpha'], fit_prior = b
best_nb.fit(x_train, y_train)
y_pred = best_nb.predict(x_test)
```

```
In [123]: print("Best Hyperparameter : ", best_param)
```

```
Best Hyperparameter : {'alpha': 0.1, 'fit_prior': True}
```

```
In [124]: acc = accuracy_score(y_test, y_pred)
acc
```

```
Out[124]: 0.9993846153846154
```

```
In [125]: print(classification_report(y_test,y_pred))
```

	precision	recall	f1-score	support
e	1.00	1.00	1.00	843
p	1.00	1.00	1.00	782
accuracy			1.00	1625
macro avg	1.00	1.00	1.00	1625
weighted avg	1.00	1.00	1.00	1625

```
In [126]: cm=confusion_matrix(y_test,y_pred)
cm
```

```
Out[126]: array([[842,  1],
                [ 0, 782]], dtype=int64)
```

```
In [127]: param_dist = {
            'alpha': uniform(0.1, 2.0), # Example: Uniform distribution for alpha
            'fit_prior':[True,False]
        }
```

```
In [128]: bayes = naive_bayes.MultinomialNB()
```

```
In [129]: x=scaler.fit_transform(x)
```

```
In [130]: from sklearn.utils.validation import check_non_negative
check_non_negative(x, "MultinomialNB (input x)")
```

```
In [131]: randomized_search = RandomizedSearchCV(bayes, param_distributions=param_dist,
randomized_search.fit(x, y)
```

```
Out[131]: RandomizedSearchCV(cv=5, estimator=MultinomialNB(),
                             param_distributions={'alpha': <scipy.stats._distn_infrastructure.rv_continuous_frozen object at 0x0000020E8513D290>,
                                                  'fit_prior': [True, False]},
                             scoring='accuracy')
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [132]: best_param = randomized_search.best_params_
print("Best Hyperparameter : ", best_param)
```

```
Best Hyperparameter : {'alpha': 0.28467359617041776, 'fit_prior': True}
```

```
In [133]: best_nb = naive_bayes.MultinomialNB(alpha = best_param['alpha'], fit_prior = b
best_nb.fit(x_train, y_train)
y_pred = best_nb.predict(x_test)
```

```
In [134]: acc = accuracy_score(y_test, y_pred)
acc
```

```
Out[134]: 0.9993846153846154
```



```
In [135]: print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
e	1.00	1.00	1.00	843
p	1.00	1.00	1.00	782
accuracy			1.00	1625
macro avg	1.00	1.00	1.00	1625
weighted avg	1.00	1.00	1.00	1625

```
In [136]: cm=confusion_matrix(y_test,y_pred)
cm
```

```
Out[136]: array([[842,  1],
                 [ 0, 782]], dtype=int64)
```

```
In [ ]:
```